




AgentDocsBench

Can coding agents actually
build from your documentation?

The verifier decides PASS or FAIL. The agent does not.





The agent can finish and still be wrong.

A page that is clear to a person can still be the wrong page for a coding agent.



THE DOCS

Pagination

The next page is `next_page`.



THE AGENT

Cursor, Claude Code, or Codex
Implemented pagination.
Done.



THE VERIFIER

pagination

FAIL

The project did not follow the docs.

Asking the agent to update the README does not check the behavior.



Don't trust "Done." Verify the implementation.

One task. The agent edits a starter. Your check decides the result.

01



Docs

The instructions the agent is allowed to read.

02



Coding agent

Cursor, Claude Code, or Codex edits the starter.

03



Project changes

Added, modified, and deleted paths are saved.

04



Verifier

Your check. Exit 0 is PASS. Anything else is FAIL.

The agent's claim is not an input.


PASS

FAIL



What happens when the documentation changes?

Rerun the same task. Swap the docs. Read the transition.



Held fixed

task, starter, agent, model, verifier

Only docs changed

the documentation directory

OBSERVED TRANSITION

FAIL became PASS on this run.

NEWLY PASSING

One run is an observation, not causal or statistical proof.



Check the behavior the docs describe.

You write the verifier. It checks the project, not the agent's summary.



Pagination

The next page follows the cursor the docs name.



Webhook verification

The signature check matches the documented scheme.



Retry behavior

Retries stop where the docs say they stop.



SDK migration

Calls follow the migration guide, not the old client.



Authentication

The request is signed the way the docs describe.



Configuration

Settings are read from the documented place.

Examples of checks you can write. Not a bundled task suite.



You bring the check. The harness runs it.

YOU BRING

Docs

What the agent may read for the task.

Starter

The project the agent is allowed to edit.

Verifier

The behavior check. Exit 0 is PASS.

THE HARNESS



Cursor, Claude Code, or Codex

The agent that edits the starter.



Project changes

Added, modified, and deleted paths are recorded.



Benchmark fingerprint

Hashes docs, starter, tasks, and verifiers.



Optional Docker isolation

Agent and verifier can run in separate containers.



Provider failures stay separate

A provider error is not recorded as a task FAIL.



Documentation experiments

Only the docs change between variants.

The same commands run in GitHub Actions.



AgentDocsBench

Test the developer experience
your AI-assisted users actually get.

OPEN SOURCE

agentdocs

```
$ agentdocs init my-benchmark --agent cursor
```

writes a small benchmark

```
$ agentdocs doctor --config agentdocs.yaml
```

checks setup, does not run an agent

```
$ agentdocs test --config agentdocs.yaml
```

the verifier decides PASS or FAIL

github.com/SaidiSouhaieb/agentdocsbench